

THE VERDICT AI

Product Requirements Document & Product Roadmap

AI-Powered Development & Greenlight Intelligence Platform

Version 1.0 | Strategic Product Definition | Target Release: Q3–Q4 2026

CONFIDENTIAL | Prepared for Product Strategy

01 — Executive Summary

The Verdict is an AI-powered development and greenlight intelligence platform designed to help studios, production companies, independent producers, and content teams make better decisions about what to develop, how to package it, how to produce it, and under what conditions it has the strongest chance of success.

The product is intentionally not positioned as a replacement for professional script readers, development executives, producers, financiers, or creative leaders. Human judgment is a foundational part of the product. The Verdict organizes evidence, surfaces patterns, models scenarios, explains risks, and accelerates analysis so that experienced professionals can spend more time exercising judgment and less time assembling fragmented information.

Core Product Thesis: Most AI tools tell users what is wrong with a script. The Verdict helps decision-makers determine what to do with it.

The V1 product combines Story Intelligence, Audience & Market Intelligence, Comparable Intelligence, Production Complexity Intelligence, Budget & Financial Intelligence, Explainable Greenlight Recommendations, and an interactive Scenario Simulator. Together these capabilities create a decision workspace that moves from script ingestion to human review to scenario modeling and ultimately to a documented development recommendation.

The long-term vision is to connect the full content lifecycle: Script → Development → Packaging → Greenlight → Production → Release → Actual Performance → Decision Replay. This creates a continuous learning system without removing the human decision-maker from the loop.

Product Promise

The Verdict does not make the decision. The Verdict makes the decision-maker better informed.

Primary Positioning

From script to greenlight. One decision intelligence platform.

Alternative category positioning: AI-Powered Development & Greenlight Intelligence; Content Decision Intelligence; Development Intelligence Platform.

02 — Product Vision & Principles

Vision

Create the intelligence layer that connects creative judgment, market evidence, financial assumptions, and production realities—helping entertainment organizations make faster, more transparent, and more defensible development decisions.

Product Principles

- Human judgment remains the final authority. AI provides evidence, analysis, options, and explanations—not autonomous greenlight decisions.
- Explainability over black-box scoring. Every recommendation must expose the primary drivers, risks, assumptions, and confidence level behind it.
- Scenario thinking over prediction theater. The product should model conditional outcomes rather than claim to predict hits with certainty.
- Decision support over script criticism. Coverage is the first input into a broader development workflow, not the end product.
- Evidence before inference. The system should distinguish verified facts, model-derived estimates, assumptions, and human opinions.
- Human override is a first-class workflow. Users can disagree, annotate, override, and document why.
- Learning through outcomes. The platform should eventually compare predictions and decisions with actual performance to improve organizational intelligence.
- Responsible AI. Uncertainty, data limitations, bias, and confidence must be visible to users.

03 — Problem & Market Opportunity

The Problem

Entertainment development decisions are fragmented across creative coverage, market research, talent conversations, budgeting, production planning, and executive judgment. A script may be evaluated by multiple teams, but the resulting information is often stored in disconnected documents, spreadsheets, emails, and meetings.

- Script coverage may identify story strengths and weaknesses without connecting them to commercial or production implications.
- Market research may identify audience demand without explaining how story or packaging choices affect that demand.
- Budgeting often happens after creative decisions have already increased production complexity.
- Talent discussions can materially change commercial potential but are difficult to model systematically.
- Distribution strategy can alter the financial profile of a project, yet is frequently treated as a later-stage decision.
- Executives must reconcile qualitative judgment with quantitative evidence under uncertainty.
- There is limited institutional memory connecting past greenlight decisions to actual outcomes.

Opportunity

The Verdict can occupy the space between traditional script coverage and enterprise greenlight decision-making. Rather than competing only on the speed or cost of AI-generated notes, it can create a differentiated workflow where creative, commercial, financial, and operational intelligence converge.

The strongest competitive advantage is not a single model or score. It is the integrated decision workflow, explainability layer, human feedback loop, scenario engine, proprietary historical data, and eventually the Decision Replay system.

04 — Target Users & Personas

Persona	Primary Need	Key Decision
Development Executive	Evaluate creative and commercial viability	Should we develop, reposition, package, or pass?
Producer	Understand feasibility and financing implications	Can we make this project at an attractive risk-adjusted profile?
Studio / Content Executive	Compare projects consistently	Which projects deserve additional development resources?
Independent Producer	Build a credible project case	What assumptions and changes improve financeability?
Production / Operations Leader	Identify physical production complexity	What will this script require to produce?
Finance / Strategy Executive	Understand downside and sensitivity	Which assumptions drive risk and potential return?
Human Script Reader / Analyst	Accelerate analysis without losing professional judgment	How can AI reduce repetitive work while improving the quality of human review?

05 — Competitive Positioning

Category Strategy

The Verdict should not compete primarily as an AI script coverage generator. That category is increasingly crowded and susceptible to commoditization. The differentiated category is Development & Greenlight Intelligence.

Traditional Coverage	Typical AI Coverage	The Verdict
Human interpretation	Automated interpretation	AI analysis + human judgment
Story-focused	Story-focused	Story + audience + market + finance + production
Static report	Static report	Interactive decision workspace
One recommendation	One score	Multiple scenarios and recommended paths

Limited quantitative modeling	Basic scoring	Explainable sensitivity and scenario analysis
Decision stored in email/docs	Report stored in app	Decision, rationale, assumptions, and overrides preserved
Little post-release learning	Little outcome feedback	Future Decision Replay and organizational learning

Core Differentiator

Most AI tools answer: “Is this a good script?” The Verdict answers: “Given this story, this audience, this talent, this budget, this production model, and this distribution strategy, what is the smartest path forward—and what would change that recommendation?”

06 — V1 Product Overview

V1 is organized as a human-centered decision pipeline:

1. Ingest the script and establish a structured project record.
2. Generate AI-assisted Story Intelligence and professional-style coverage.
3. Analyze audience, market positioning, and comparable titles.
4. Estimate production complexity and major cost drivers.
5. Build a financial and commercial baseline using explicit assumptions.
6. Generate an explainable preliminary Verdict with confidence and risks.
7. Invite human review, annotation, correction, and override.
8. Run alternative scenarios through the Scenario Simulator.
9. Compare scenarios and document the preferred development path.
10. Export an executive-ready Decision Report.

Core User Journey

Upload Script → AI Analysis → Human Review → Market & Production Intelligence → Financial Baseline → Preliminary Verdict → Scenario Simulation → Human Decision → Decision Record

07 — Feature Inventory & Priorities

ID	Feature	Description	Priority	Primary User
ST-001	Script Ingestion	Secure PDF/script upload, parsing, scene segmentation, metadata extraction	P0	All
ST-002	Story Intelligence	Structure, premise, character, pacing, conflict, themes, originality, strengths and risks	P0	Development
ST-003	Professional Coverage Workspace	AI-generated first-pass coverage with editable human review sections	P0	Readers / Development
ST-004	Audience Intelligence	Audience profile, demographic hypotheses, emotional appeal, audience risks	P0	Development / Marketing
MI-001	Market Positioning	Genre, format, market fit, saturation, strategic positioning	P0	Development
MI-002	Comparable Intelligence	Evidence-based comparable titles with	P0	Development

		similarity dimensions and rationale		
PI-001	Production Complexity	Locations, pages, characters, stunts, VFX, builds, complexity drivers	P0	Producer / Production
FI-001	Budget Baseline	Indicative budget range with assumptions and confidence	P0	Producer / Finance
GI-001	Explainable Verdict	Greenlight path, top drivers, risks, assumptions, confidence	P0	Executives
SS-001	Scenario Variables	Budget, talent strength, marketing, distribution, location and production assumptions	P0	Executives
SS-002	Real-Time Scenario Inference	Recalculate decision indicators after scenario changes	P0	Executives
SS-003	Scenario Comparison	Save and compare multiple development paths	P1	Executives
DW-001	Human Review & Override	Approve, edit, challenge, annotate, and override AI findings	P0	Human reviewers
DW-002	Decision Record	Capture final decision, rationale, assumptions, and contributors	P1	Executives
DW-003	Team Collaboration	Comments, assignments, review status, version history	P1	Teams
PR-001	Production Digital Twin	Translate script complexity into production requirements and planning assumptions	P1	Production
EX-001	Executive Report	Generate boardroom-ready report with evidence and scenario comparisons	P0	Executives
DR-001	Decision Replay Foundation	Store predictions, assumptions, overrides and final outcomes for future learning	P2	Enterprise

08 — P0 Functional Requirements

8.1 Script Ingestion & Project Setup

- Accept screenplay/script PDFs and supported structured formats.
- Extract scenes, pages, characters, locations, INT/EXT, day/night, dialogue, action, and production indicators.
- Create a persistent Project ID and versioned source document.
- Flag parsing uncertainty for human review.
- Preserve source references so findings can be traced back to pages/scenes.

8.2 Story Intelligence

- Premise and logline analysis.
- Narrative structure and pacing analysis.
- Character arc and relationship analysis.
- Conflict, stakes, themes, tone, and genre analysis.
- Strengths and weaknesses with page-level evidence.
- Potential audience appeal and barriers.
- Revision opportunities prioritized by likely impact.
- Clear distinction between observed evidence and AI interpretation.

8.3 Professional Coverage Workspace

The AI generates a structured first pass, but the human reader remains an active participant. The workspace should include editable sections for logline, synopsis, comments, strengths, weaknesses, marketability, production considerations, recommendation, and reader notes.

- Human can accept, edit, reject, or rewrite AI-generated findings.
- Human can add comments not detected by AI.
- Human can mark AI findings as inaccurate or unsupported.
- AI can explain the evidence supporting a finding.
- Final coverage is clearly labeled as human-reviewed.

8.4 Audience & Market Intelligence

- Identify primary and secondary audience hypotheses.
- Estimate audience breadth versus niche appeal.
- Identify genre and format demand signals.
- Assess competitive saturation.
- Evaluate theatrical, streaming, premium cable, and other potential positioning.
- Surface audience risks and unmet opportunities.
- Allow human users to challenge or replace AI audience assumptions.

8.5 Comparable Intelligence

- Identify relevant comparable titles.
- Score similarity across narrative, audience, genre, budget, production model, and distribution—not just premise.
- Explain why each comparable matters.
- Separate creative comparables from commercial comparables.
- Show data freshness and source confidence.
- Allow human users to add or remove comparables.

8.6 Production Complexity Intelligence

- Estimate shooting days from script complexity and page count.
- Identify location count and recurring locations.
- Flag heavy builds and construction requirements.
- Identify stunts, action, VFX, SFX, crowd, water, vehicles, and other complexity drivers.
- Identify potential virtual production candidates.
- Estimate stage footprint requirements where sufficient data exists.
- Expose assumptions rather than presenting estimates as final production budgets.

8.7 Budget & Financial Intelligence

- Generate indicative budget range rather than false precision.
- Identify major cost drivers.
- Model budget sensitivity.
- Allow users to enter known talent, location, incentive, and production assumptions.
- Show confidence intervals and uncertainty.
- Separate modeled assumptions from verified financial data.
- Support scenario-level comparisons.

8.8 Explainable Verdict

The Verdict should use a decision path rather than a single black-box score. Recommended statuses:

- GREENLIGHT — Evidence supports proceeding to packaging or production planning.
- DEVELOP — Strong potential, but specific creative or commercial issues should be addressed.
- REPOSITION — Opportunity exists, but budget, format, audience, or distribution strategy should change.
- PACKAGE — The project is highly dependent on talent, financing, or distribution attachment.
- PASS — Current risk materially exceeds potential under modeled assumptions.

Every Verdict must display: primary positive drivers, primary risk drivers, assumptions, confidence level, data limitations, and the conditions most likely to change the recommendation.

09 — Scenario Simulator

The Scenario Simulator is the central V1 differentiator. It transforms The Verdict from a static report into an interactive decision system.

Scenario Inputs

- Production budget
- Marketing spend
- Talent strength / bankability
- Production location
- Theatrical vs. streaming distribution
- Release scale
- VFX / virtual production strategy
- Production complexity assumptions
- Target audience positioning

Scenario Outputs

- Risk-adjusted opportunity indicator
- Commercial potential range
- Budget sensitivity
- Production complexity
- Audience reach potential
- Distribution fit
- Confidence level
- Top factors driving change
- Recommended development path

Example

Scenario A — Tentpole: high budget, premium talent, wide release. Scenario B — Mid-budget: recognizable talent, controlled production, moderate marketing. Scenario C — Streaming Original: lower budget, emerging talent, platform-first release. The system should identify which scenario produces the strongest risk-adjusted profile and explain why.

What Would Change the Verdict?

The system should explicitly identify the smallest number of changes most likely to materially alter the recommendation, such as reducing budget, attaching talent, changing distribution strategy, reducing VFX exposure, strengthening a character arc, or repositioning the target audience.

10 — Human-in-the-Loop Decision Workspace

Human contribution is a core product capability. The Verdict should be designed around a collaborative workflow where AI accelerates analysis and humans contribute context, taste, experience, relationships, institutional knowledge, and final judgment.

- Human Review Queue: prioritize AI findings requiring review.
- Evidence View: trace findings to script pages, scenes, or external data.
- Challenge Finding: user can disagree with an AI assessment and provide rationale.
- Override: human can replace an AI recommendation.
- Confidence Adjustment: authorized users can adjust confidence when proprietary knowledge changes the assessment.
- Decision Notes: preserve why the final decision differed from the AI recommendation.
- Attribution: record human contributors and roles.
- Versioning: preserve AI baseline, human revisions, and final decision.

The system should never obscure the difference between AI-generated analysis and human-authored judgment. The final report should clearly identify both.

11 — AI Architecture & Intelligence Systems

Architectural Principle

The LLM should function as a reasoning and orchestration layer—not merely a text generator. Structured extraction, deterministic calculations, retrieval, statistical models, and human inputs should feed the reasoning process.

Layer	Purpose	Approach
Script Understanding	Parse and structure screenplay	Document parsing + NLP + structured extraction
Story Intelligence	Analyze narrative and characters	LLM reasoning with source citations
Market Intelligence	Assess market fit and competition	RAG + structured market datasets
Comparable Engine	Find and explain relevant comps	Vector + metadata retrieval + LLM reasoning
Production Intelligence	Estimate complexity and requirements	Rule-based extraction + NLP + production heuristics
Financial Intelligence	Model ranges and sensitivities	Deterministic formulas + statistical models
Scenario Engine	Recalculate conditional outcomes	Scenario orchestration + sensitivity analysis
Verdict Layer	Synthesize evidence	LLM reasoning with confidence and explainability
Human Review Layer	Capture professional judgment	Structured annotations, overrides, and decision records
Future Decision Replay	Learn from outcomes	Historical predictions vs. actual performance

AI Guardrails

- No unsupported factual claims presented as facts.
- Every major conclusion includes evidence or an explicit assumption.
- Confidence is shown at the component level where practical.
- Model uncertainty is surfaced to users.
- AI cannot silently override human decisions.
- Sensitive or proprietary project data is isolated and access-controlled.
- Training and model-improvement use of customer data requires explicit policy and consent.

12 — Data Strategy

Core Data Domains

- Script and screenplay data
- Scene and production metadata
- Historical film and television performance
- Genre and audience data

- Comparable title metadata
- Budget ranges and production economics
- Distribution and release strategy
- Talent and packaging information where legally and ethically available
- Human coverage and decision annotations
- Actual post-release performance

Data Moat

The long-term moat should be the accumulation of structured decision data: what humans believed, what AI inferred, what assumptions were made, what decisions were taken, and what actually happened. Over time, The Verdict can become better at identifying which factors matter most for specific genres, formats, budgets, platforms, and organizations.

13 — Technical Architecture

Recommended initial architecture follows the same serverless-first philosophy proven in the AirDate.tv product strategy: managed compute, API orchestration, object storage, scalable database services, model APIs, observability, and strong security controls.

Component	Recommended Role
Frontend	React / Next.js or equivalent decision workspace
Authentication	AWS Cognito or enterprise SSO
API	API Gateway + Lambda or equivalent managed API layer
Script Storage	Amazon S3 with encryption and versioning
Structured Data	DynamoDB / PostgreSQL depending on relational needs
Search / Retrieval	OpenSearch with vector search or equivalent
LLM	Amazon Bedrock model routing with model abstraction
ML	SageMaker or managed model endpoints for structured predictive tasks
Workflow	Step Functions / EventBridge for asynchronous analysis
Reporting	Server-side PDF/DOCX generation
Observability	CloudWatch, CloudTrail, centralized logging
Security	WAF, GuardDuty, IAM least privilege, encryption, secrets management
Analytics	Product analytics plus model evaluation telemetry

14 — Security, Privacy & Trust

- Encryption in transit and at rest.
- Tenant isolation for enterprise customers.
- Role-based access control.
- Audit logging for document access and decision changes.
- Versioned source documents and analysis outputs.
- Configurable data retention.
- No customer script used for model training without explicit permission.
- Secure deletion workflows.
- Human review and AI provenance metadata.
- Disaster recovery and business continuity planning.

Because scripts and development materials can be highly confidential, security and privacy should be treated as product features, not post-launch infrastructure work.

15 — Success Metrics & KPIs

Metric	Definition	V1 Target / Direction
--------	------------	-----------------------

Decision Velocity	Time from script submission to documented development recommendation	Reduce materially versus baseline
Human Review Efficiency	Time saved in first-pass analysis without reducing review quality	20–40% improvement target
Evidence Traceability	% of major findings traceable to source evidence or assumptions	>90%
Human Override Rate	% of AI findings changed by human reviewers	Measure, don't optimize downward
Decision Confidence	User-rated confidence in final decision process	Increase over baseline
Scenario Adoption	% of projects using at least two scenarios	>60% of active projects
Recommendation Usefulness	Human rating of AI recommendation usefulness	>80% positive
Prediction Accuracy	Variance between modeled expectations and actual outcomes	Track by segment
Outcome Learning	% of completed projects with actual outcomes captured	Build toward comprehensive coverage
Retention	Teams returning for multiple projects	Primary enterprise retention KPI

16 — V1 Roadmap

Phase 0 — Foundation & Validation | 0–8 Weeks

- Finalize product positioning and ICP.
- Conduct 10–15 interviews with development executives, producers, script readers, and independent producers.
- Validate the human-in-the-loop workflow.
- Define data taxonomy and confidence framework.
- Create evaluation dataset of representative scripts and human coverage.
- Establish baseline AI quality metrics.
- Prototype executive Verdict report.
- Define security and data-retention requirements.

Phase 1 — V1 MVP: Story-to-Verdict | Months 2–4

- Script ingestion and project workspace.
- Story Intelligence.
- Professional Coverage Workspace.
- Audience & Market Intelligence.
- Comparable Intelligence.
- Explainable Verdict.
- Human review, annotation, and override.
- Executive Decision Report.
- Basic authentication and project history.

Phase 2 — V1 Decision Intelligence | Months 4–6

- Budget baseline.
- Production Complexity Intelligence.
- Scenario Simulator.
- Scenario variables and real-time inference.
- Scenario comparison.
- What Would Change the Verdict?
- Decision Record.
- Team collaboration and review status.
- Enterprise security hardening.

Phase 3 — V1.5 Production Intelligence | Months 6–9

- Production Digital Twin foundation.
- Stage-to-page estimator.
- Location and set footprint analysis.
- Construction and prep lead-time indicators.
- VFX and virtual production flags.
- More granular budget sensitivity.
- Location and incentive scenario modeling.
- Pilot with production operations users.

Phase 4 — V2 Organizational Intelligence | Months 9–15

- Decision Replay foundation.
- Prediction vs. actual performance tracking.
- Historical project benchmarking.
- Organization-specific decision patterns.
- Portfolio dashboard.
- Greenlight pipeline.
- Advanced collaboration and approvals.
- Private enterprise data integrations.

Phase 5 — V3 Strategic Intelligence | Months 15–24

- Enterprise-grade organizational learning.
- Cross-project pattern detection.
- Advanced risk-adjusted portfolio optimization.
- Integration with production management and finance systems.
- Studio-specific intelligence models.
- API and data products.
- Industry benchmarking where legally and ethically permitted.

17 — Roadmap at a Glance

Release	Theme	Primary Outcome
V0	Validation	Prove human-centered workflow and decision value
V1.0	Story-to-Verdict	Turn a script into an explainable development recommendation
V1.1	Decision Intelligence	Model scenarios and identify the strongest path forward
V1.5	Production Intelligence	Connect creative decisions to physical production realities
V2.0	Organizational Intelligence	Learn from decisions and outcomes
V3.0	Strategic Intelligence	Optimize portfolios and enterprise development strategy

18 — Product Moat & Defensibility

- Integrated workflow across creative, market, financial, and production intelligence.
- Human-in-the-loop decision history.
- Proprietary decision and outcome dataset.
- Explainable scenario modeling.
- Organization-specific benchmarks.
- Production Digital Twin capabilities.
- Decision Replay feedback loop.
- Enterprise trust, security, and auditability.
- Deep integration into the development workflow creates high switching costs.

The long-term defensibility of The Verdict should come from the system of intelligence around decisions—not from any single LLM. Models will improve and become cheaper. The proprietary combination of structured project data, human judgments, scenario assumptions, organizational decision patterns, and actual outcomes is substantially harder to replicate.

19 — Go-to-Market Strategy

Initial Beachhead

Start with independent producers, boutique production companies, and development teams that need professional-quality analysis but lack the resources of major studios. These users can provide rapid feedback and generate diverse project data.

Enterprise Expansion

Use validated V1 workflows to expand into studio development groups, content strategy teams, production operations, and finance.

Commercial Model

- Professional: per-project or monthly subscription for independent producers and consultants.
- Team: collaborative workspace with project history and shared decision records.
- Enterprise: annual contract with SSO, private data controls, custom benchmarks, and integrations.
- Future Data/API: aggregated intelligence products where data rights and privacy permit.

20 — Risks & Mitigations

Risk	Mitigation
Users treat AI recommendations as facts	Confidence intervals, evidence links, assumptions, human approval
AI generates generic coverage	Benchmark against expert readers and continuously evaluate quality
Predictive claims damage credibility	Position as decision support, not hit prediction
Human users feel displaced	Make human review, override, and attribution first-class features
Sensitive scripts create security concerns	Enterprise security, encryption, access controls, retention policies
Model bias influences recommendations	Bias evaluation, diverse evaluation sets, human review
Data quality limits commercial analysis	Source confidence, freshness indicators, explicit uncertainty
Too broad for V1	Sequence Story-to-Verdict first, Scenario Intelligence second, Production Intelligence third
Complex enterprise sales cycle	Begin with producer/development workflows and prove measurable decision velocity

21 — Recommended V1 Definition of Done

- A user can upload a script and receive a structured analysis with source evidence.
- A human reader can edit, challenge, and override AI findings.
- The system can produce audience, market, and comparable intelligence.
- The system can estimate production complexity and budget ranges with explicit assumptions.
- The system generates an explainable preliminary Verdict.
- A user can manipulate core scenario variables.
- The system can compare at least three scenarios.
- The system identifies what changes would most likely alter the recommendation.
- A human can document the final decision and rationale.
- The system preserves AI findings, human edits, and final decisions.
- The system generates an executive-ready report.
- Security and privacy controls are appropriate for confidential development materials.

22 — Final Product Positioning

The Verdict is not an AI script reader. It is not an automated greenlight committee. It is not a black-box hit predictor.

The Verdict is an AI-powered development intelligence platform that brings together creative analysis, market intelligence, financial modeling, production intelligence, and scenario simulation—so experienced entertainment professionals can make faster, more informed, and more defensible decisions.

The fundamental product philosophy is simple:

AI accelerates the analysis. Humans provide the judgment. The Verdict connects the two.

Appendix A — Suggested V1 Executive Report Structure

11. Project Overview
12. Human Coverage Summary
13. Story Intelligence
14. Audience & Market Intelligence
15. Comparable Intelligence
16. Production Complexity
17. Budget & Financial Baseline
18. Preliminary Verdict
19. Top Positive Drivers
20. Top Risk Drivers
21. Confidence & Data Limitations
22. What Would Change the Verdict?
23. Scenario A / B / C Comparison
24. Human Reviewer Notes
25. Final Decision & Rationale
26. Next Recommended Actions

Appendix B — North Star

The North Star for The Verdict is not prediction accuracy alone. It is better decision quality.

A successful product should help a development team reach a decision faster, understand why it reached that decision, identify what assumptions matter most, preserve human expertise, and learn from the eventual outcome.

That is the foundation for building a durable content decision intelligence platform.